Rule-Based Anomaly Detection

- 1) Problem Presentation
 - 2) Demonstration
 - 3) Opinion

Lisa Robert
Chirine Makhlouf
Alison Lopes-Sousa
Anissa Cedrati
Kevin Fonseca
Jordan Domingues

1) Problem presentation

The problem: hackers trying to gain access to a website using a brute-force attack.

The problem presented in this chapter is that we are facing the menace of hackers entering our website via a **brute-force attack**. They are using this technique which consists in using **many combinations of usernames and passwords until a combination works**.

1) Problem presentation

To address this problem we will be using **anomaly detection** which is a process used to identify events, or observations that are different from what's considered normal or typical in a dataset. When a data point or event is identified as an anomaly it means that it **deviates from the usual pattern** or **behavior of the dataset**.

By noticing these anomalies we will be able to recognize a suspicious activity and stop it before the hackers get into the website.

To test our process we will be simulating attacks

Complete the following exercises to practice the concepts covered in this chapter:

- 1. Run the simulation for December 2018 into new log files without making the user base again. Be sure to run python3 simulate.py -h to review the command-line arguments. Set the seed to 27. This data will be used for the remaining exercises.
- 2. Find the number of unique usernames, attempts, successes, and failures, as well as the success/failure rates per IP address, using the data simulated from exercise 1.
- 3. Create two subplots with failures versus attempts on the left, and failure rate versus distinct usernames on the right. Draw decision boundaries for the resulting plots. Be sure to color each data point by whether or not it is a hacker IP address.
- 4. Build a rule-based criteria using the percentage difference from the median that flags an IP address if the failures and attempts are both five times their respective medians, or if the distinct usernames count is five times its median. Be sure to use a one-hour window. Remember to use the get_baselines() function to calculate the metrics needed for the baselines.
- 5. Calculate metrics to evaluate how well these rules performed using the evaluate() and classification_stats() functions from this chapter.

2) Demonstration

EXERCISE 1

Run the simulation for December 2018 into new log files without making the user base again. Be sure to run python3 simulate.py -h to review the command-line arguments. Set the seed to 27. This data will be used for the remaining exercises.

[INFO] [simulate.py] Simulating 31.0 days... [INFO] [simulate.py] Saving logs

[INFO] [simulate.py] All done!

```
Entrée [7]:

1 %matplotlib inline
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import pandas as pd
5 import seaborn as sns
```

Alison LOPES SOUSA

2) Demonstration

EXERCISE 2

Find the number of unique usernames, attempts, successes, and failures, as well as the success/failure rates per IP address, using the data simulated from exercise 1.

Alison LOPES SOUSA

2) Demonstration

EXERCISE 2

Find the number of unique usernames, attempts, successes, and failures, as well as the success/failure rates per IP address, using the data simulated from exercise 1.



```
Édition
               Affichage
                         Langage
 1 datetime, source_ip, username, success, failure_reason
 2 2018-12-01 00:23:11.097544,75.161.66.106,olopez,True,
 3 2018-12-01 00:41:49.098371,95.227.178.199,olopez,False,error wrong password
 4 2018-12-01 00:41:50.098371,95.227.178.199,olopez,True,
 5 2018-12-01 01:02:59.347432,87.170.218.31,bbrown,True,
 6 2018-12-01 01:04:21.404189,87.170.218.31,bbrown,True,
 7 2018-12-01 01:10:11.031665,51.74.112.23,bbrown,True,
 8 2018-12-01 01:19:07.580804,75.161.66.106,olopez,True,
 9 2018-12-01 01:52:32.449943,95.227.178.199,olopez,True,
10 2018-12-01 02:16:13.543327,218.251.199.112,olopez,True,
11 2018-12-01 02:17:09.206386,198.92.176.42,xkim,True,
12 2018-12-01 02:20:00.201812,198.92.176.42,xkim,True,
13 2018-12-01 02:51:10.101921,95.227.178.199,olopez,True,
14 2018-12-01 02:54:43.317021,95.227.178.199,olopez,True,
15 2018-12-01 02:59:48.825331,22.64.52.205,bsmith,False,error wrong password
16 2018-12-01 02:59:49.825331,22.64.52.205,bsmith,False,error_wrong_password
17 2018-12-01 02:59:50.825331,22.64.52.205, ibeown, False, error wrong username
18 2018-12-01 02:59:51.825331,22.64.52.205, ibeown, False, error wrong username
19 2018-12-01 02:59:52.825331,22.64.52.205,gloez,False,error wrong username
   2018-12-01 02:59:53.825331,22.64.52.205,glopez,False,error wrong password
```

Alison LOPES SOUSA

2) Demonstration

EXERCISE 2

Find the number of unique usernames, attempts, successes, and failures, as well as the success/failure rates per IP address, using the data simulated from exercise 1.

:		source_ip	username	success	failures	attempts	success_rate	failure_rate
	0	1.138.149.116	1	14	0	14	1.000000	0.000000
	1	1.54.45.23	49	17	68	85	0.200000	0.800000
	2	100.43.18.36	1	7	0	7	1.000000	0.000000
	3	101.113.31.197	1	10	1	11	0.909091	0.090909
	4	101.154.143.93	1	30	5	35	0.857143	0.142857

EXERCISE 3

Create two subplots with failures versus attempts on the left, and failure rate versus distinct usernames on the right. Draw decision boundaries for the resulting plots. Be sure to color each data point by whether or not it is a hacker IP address.

```
Entrée [9]:
             1 is_attack_ip = log_aggs.source_ip.isin(
                    pd.read_csv('dec_2018_attacks.csv').source_ip
                fig, axes = plt.subplots(1, 2, figsize=(10, 3))
                for ax, (x, y) in zip(axes, (('attempts', 'failures'), ('username', 'failure rate'))):
                    ax = sns.scatterplot(
                        x=log_aggs[x],
                        y=log_aggs[y],
                        hue=is_attack_ip,
            12
                        ax=ax
            13
                    ax.set_title(f'{y.title()} vs. {x.title()}')
            14
            15
            16 # boundaries
                axes[0].plot([0, 80], [80, 0], 'r--')
            18 axes[1].axhline(0.5, color='red', linestyle='--')
```

Out[9]: <matplotlib.lines.Line2D at 0x7f949d3c0370>

EXERCISE 3

Create two subplots with failures versus attempts on the left, and failure rate versus distinct usernames on the right. Draw decision boundaries for the resulting plots. Be sure to color each data point by whether or not it is a hacker IP address.



EXERCISE 3

Create two subplots with failures versus attempts on the left, and failure rate versus distinct usernames on the right. Draw decision boundaries for the resulting plots. Be sure to color each data point by whether or not it is a hacker IP address.

```
Entrée [9]:
             1 is_attack_ip = log_aggs.source_ip.isin(
                    pd.read_csv('dec_2018_attacks.csv').source_ip
                fig, axes = plt.subplots(1, 2, figsize=(10, 3))
                for ax, (x, y) in zip(axes, (('attempts', 'failures'), ('username', 'failure rate'))):
                    ax = sns.scatterplot(
                        x=log_aggs[x],
                        y=log_aggs[y],
                        hue=is_attack_ip,
            12
                        ax=ax
            13
                    ax.set_title(f'{y.title()} vs. {x.title()}')
            14
            15
            16 # boundaries
                axes[0].plot([0, 80], [80, 0], 'r--')
            18 axes[1].axhline(0.5, color='red', linestyle='--')
```

Out[9]: <matplotlib.lines.Line2D at 0x7f949d3c0370>

EXERCISE 3

Create two subplots with failures versus attempts on the left, and failure rate versus distinct usernames on the right. Draw decision boundaries for the resulting plots. Be sure to color each data point by whether or not it is a hacker IP address.

	source_ip	username	success	failures	attempts	success_rate	failure_rate
0	1.138.149.116	1	14	0	14	1.000000	0.000000
1	1.54.45.23	49	17	68	85	0.200000	0.800000
2	100.43.18.36	1	7	0	7	1.000000	0.000000
3	101.113.31.197	1	10	1	11	0.909091	0.090909
4	101.154.143.93	1	30	5	35	0.857143	0.142857

EXERCISE 3

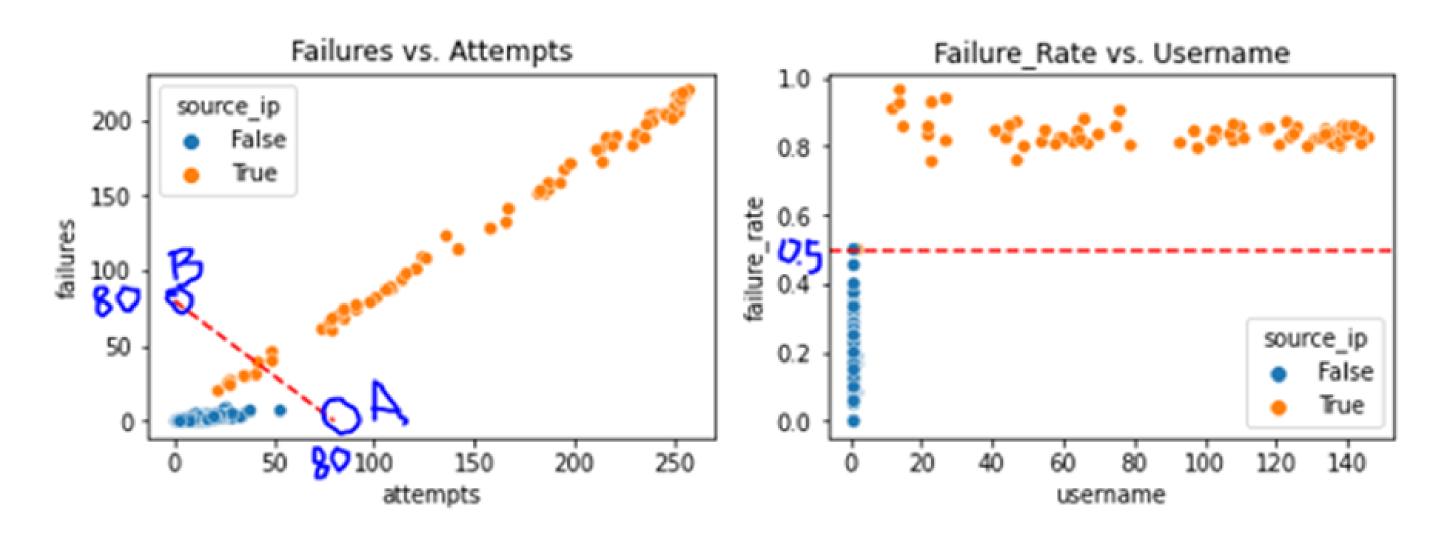
Create two subplots with failures versus attempts on the left, and failure rate versus distinct usernames on the right. Draw decision boundaries for the resulting plots. Be sure to color each data point by whether or not it is a hacker IP address.

```
Entrée [9]:
             1 is_attack_ip = log_aggs.source_ip.isin(
                    pd.read_csv('dec_2018_attacks.csv').source_ip
                fig, axes = plt.subplots(1, 2, figsize=(10, 3))
                for ax, (x, y) in zip(axes, (('attempts', 'failures'), ('username', 'failure rate'))):
                    ax = sns.scatterplot(
                        x=log_aggs[x],
                        y=log_aggs[y],
                        hue=is_attack_ip,
            12
                        ax=ax
            13
                    ax.set_title(f'{y.title()} vs. {x.title()}')
            14
            15
            16 # boundaries
                axes[0].plot([0, 80], [80, 0], 'r--')
            18 axes[1].axhline(0.5, color='red', linestyle='--')
```

Out[9]: <matplotlib.lines.Line2D at 0x7f949d3c0370>

EXERCISE 3

Create two subplots with failures versus attempts on the left, and failure rate versus distinct usernames on the right. Draw decision boundaries for the resulting plots. Be sure to color each data point by whether or not it is a hacker IP address.



EXERCISE 4

Build a rule-based criteria using the percentage difference from the median that flags an IP address if the failures and attempts are both five times their respective medians, or if the distinct usernames count is five times its median. Be sure to use a one-hour window. Remember to use the get_baselines() function to calculate the metrics needed for the baselines.

Our to 1	r 1 A 1	-
MALL I	1 1 1/1 1	
		_

H		source_ip	datetime	username	success	failures	attempts	success_rate	failure_rate
	0	1.138.149.116	2018-12-11 13:00:00	1	3	0	3	1.0	0.0
	1	1.138.149.116	2018-12-11 14:00:00	1	4	0	4	1.0	0.0
	2	1.138.149.116	2018-12-15 20:00:00	1	2	0	2	1.0	0.0
	3	1.138.149.116	2018-12-17 23:00:00	1	2	0	2	1.0	0.0
	4	1.138.149.116	2018-12-18 00:00:00	1	1	0	1	1.0	0.0

EXERCISE 4

Build a rule-based criteria using the percentage difference from the median that flags an IP address if the failures and attempts are both five times their respective medians, or if the distinct usernames count is five times its median. Be sure to use a one-hour window. Remember to use the get_baselines() function to calculate the metrics needed for the baselines.

```
def get baselines(hourly ip logs, func, *args, **kwargs):
Entrée [11]:
                     Calculate hourly bootstrapped statistic per column.
                     Parameters:
                         - hourly ip logs: Data to sample from.
                         - func: Statistic to calculate.
                         - args: Additional positional arguments for `func`
                         - kwargs: Additional keyword arguments for `func`
              10
                     Returns:
              12
                          pandas.DataFrame of hourly bootstrapped statistics
              13
              14
                     if isinstance(func, str):
                         func = getattr(pd.DataFrame, func)
              15
              16
                     return hourly_ip_logs.assign(
              17
              18
                         hour=lambda x: x.datetime.dt.hour
                     ).groupby('hour').apply(
              19
                         lambda x: x.sample(10, random_state=0, replace=True).pipe(func, *args, **kwargs, numeric_only=True)
              20
              21
```

EXERCISE 4

Build a rule-based criteria using the percentage difference from the median that flags an IP address if the failures and attempts are both five times their respective medians, or if the distinct usernames count is five times its median. Be sure to use a one-hour window. Remember to use the get_baselines() function to calculate the metrics needed for the baselines.

2) Demonstration

EXERCISE 5

Calculate metrics to evaluate how well these rules performed using the evaluate() and classification_stats() functions from this chapter.

```
Entrée [14]:
               1 def evaluate(alerted ips, attack ips, log ips):
                     Calculate true positives (TP), false positives (FP),
                     true negatives (TN), and false negatives (FN) for
                     IP addresses flagged as suspicious.
                      Parameters:
                         - alerted ips: `pandas.Series` of flagged IP addresses
                         - attack ips: `pandas.Series` of attacker IP addresses
              10
                          - log ips: `pandas.Series` of all IP addresses seen
              11
              12
                     Returns:
                          Tuple of form (TP, FP, TN, FN)
              14
              15
                     tp = alerted ips.isin(attack ips).sum()
                     tn = np.invert(np.isin(log_ips[~log_ips.isin(alerted_ips)].unique(), attack_ips)).sum()
              17
                     fp = np.invert(
              18
                          np.isin(log ips[log ips.isin(alerted ips)].unique(), attack ips)
              19
                     ).sum()
                     fn = np.invert(
              21
                          np.isin(log ips[log ips.isin(attack ips)].unique(), alerted ips)
              22
                      ).sum()
              23
                     return tp, fp, tn, fn
```

2) Demonstration

EXERCISE 5

Calculate metrics to evaluate how well these rules performed using the evaluate() and classification_stats() functions from this chapter.

```
1 # make this easier to call
Entrée [15]:
               2 from functools import partial
               3 scores = partial(
                     evaluate,
                     attack_ips=pd.read_csv('dec_2018_attacks.csv').source_ip,
                     log_ips=dec_log.source_ip.drop_duplicates()
Entrée [16]:
               1 def classification_stats(tp, fp, tn, fn):
                     """Calculate metrics"""
                     return {
                         'FPR': fp / (fp + tn),
                       'FDR': fp / (fp + tp),
                         'FNR': fn / (fn + tp),
                         'FOR': fn / (fn + tn)
              1 classification_stats(*scores(flagged_ips))
Entrée [17]:
  Out[17]: {'FPR': 0.07003891050583658,
            'FDR': 0.18947368421052632,
            'FNR': 0.01282051282051282,
            'FOR': 0.00416666666666667}
```

2) Demonstration

EXERCISE 5

Calculate metrics to evaluate how well these rules performed using the evaluate() and classification_stats() functions from this chapter.

```
File Edit View Language

1 start,end, source_ip
2 2018-12-01 02:59:47.825331,2018-12-01 03:03:28.825331,22.64.52.205
3 2018-12-01 09:48:44.833872,2018-12-01 09:50:30.833872,88.55.2.156
4 2018-12-01 18:17:17.328128,2018-12-01 18:19:39.328128,165.230.10.86
5 2018-12-02 03:01:47.695193,2018-12-02 03:04:49.605193,120.238.96.119
6 2018-12-02 04:40:40.605517,2018-12-02 04:44:57.605517,77.195.252.66.
7 2018-12-03 03:42:50.812476,2018-12-03 03:46:50.812476,137.146.98.203
8 2018-12-03 15:38:02.757783,2018-12-03 07:02:38.130102,104.14.16.33.
9 2018-12-03 15:38:02.757783,2018-12-03 15:38:04.757783,203.100.15.211
10 2018-12-04 08:54:35.511358,2018-12-04 08:58:46.511358,161.0.204.113
11 2018-12-05 11:19:54.655905,2018-12-05 11:20:22.655905,206.196.247.152
12 2018-12-06 21:46:52.985600,2018-12-06 21:48:14.985600,132.9.183.215
14 2018-12-07 00:30:51.023973,2018-12-07 00:53:42.445482,118.73.99.180
```

3) Opinion

Intellectual aspect

- Difficulties understanding the exercice.
- Confusing instructions & difficult to follow.
- Programming & coding are unknown areas for us.
- No initiation& training in these fields.
- Poor knowledge.
- Connection with our training / Master?

Jordan DOMINGUES Kévin FONSECA

Technical aspect

- Problems of exploitation.
- Different operating systems.
- Not the same ease of use.
- Exercice should have been shared directly and simultaneously.
- Problems ordering & recording on the program.

THANK YOU FOR YOUR ATTENTION